Try again once you are ready.

Required to pass: 80% or higher

You can retake this quiz up to 3 times every 8 hours.

Back to Week 5

Retake

0 / 1 point

1. You are training a three layer neural network and would like to use backpropagation to compute the gradient of the cost function. In the backpropagation algorithm, one of the steps is to update

$$\Delta_{ij}^{(2)} := \Delta_{ij}^{(2)} + \delta_i^{(3)} * (a^{(2)})_j$$

for every i, j. Which of the following is a correct vectorization of this step?

- $\Delta^{(2)} := \Delta^{(2)} + (a^{(2)})^T * \delta^{(3)}$

This should not be selected

This choice adds $(a^{(2)})^T * \delta^{(3)}$ to $\Delta^{(2)}$, but this value is a scalar and we need to add a matrix to update every entry at once.

- $\Delta^{(2)} := \Delta^{(2)} + (a^{(3)})^T * \delta^{(2)}$
- $\Delta^{(2)} := \Delta^{(2)} + \delta^{(3)} * (a^{(2)})^T$
- $\Delta^{(2)} := \Delta^{(2)} + \delta^{(3)} * (a^{(3)})^T$

1 / 1 point

2. Suppose **Theta1** is a 5x3 matrix, and **Theta2** is a 4x6 matrix. You set $\texttt{thetaVec} = [\texttt{Theta1}(:); \texttt{Theta2}(:)]$. Which of the following correctly recovers **Theta2**?

- $\texttt{reshape}(\texttt{thetaVec}(16:39), 4, 6)$

  Correct
  This choice is correct, since **Theta1** has 15 elements, so **Theta2** begins at index 16 and ends at index 16 + 24 - 1 = 39.

- $\texttt{reshape}(\texttt{thetaVec}(15:38), 4, 6)$
- $\texttt{reshape}(\texttt{thetaVec}(16:24), 4, 6)$
- $\texttt{reshape}(\texttt{thetaVec}(15:39), 4, 6)$
- $\texttt{reshape}(\texttt{thetaVec}(16:39), 6, 4)$

---

1 / 1 point

3. Let $J(\theta) = 3\theta^4 + 4$. Let $\theta = 1$, and $\epsilon = 0.01$. Use the formula $\frac{J(\theta+\epsilon)-J(\theta-\epsilon)}{2\epsilon}$ to numerically compute an approximation to the derivative at $\theta = 1$. What value do you get? (When $\theta = 1$, the true/exact derivative is $\frac{dJ(\theta)}{d\theta} = 12$.)

- 12
- 6
- 12.0012

  Correct

We compute $\frac{(3(1.01)^4+4)-(3(0.99)^4+4)}{2(0.01)} = 12.0012.$

- [ ] 11.9988

✔ 1 / 1 point

4. Which of the following statements are true? Check all that apply.

- [ ] Computing the gradient of the cost function in a neural network has the same efficiency when we use backpropagation or when we numerically compute it using the method of gradient checking.

Un-selected is correct

- [ ] Gradient checking is useful if we are using one of the advanced optimization methods (such as in fminunc) as our optimization algorithm. However, it serves little purpose if we are using gradient descent.

Un-selected is correct

- [ ] Using gradient checking can help verify if one's implementation of backpropagation is bug-free.

Correct

If the gradient computed by backpropagation is the same as one computed numerically with gradient checking, this is very strong evidence that you have a correct implementation of backpropagation.

- [ ] For computational efficiency, after we have performed gradient checking to

verify that our backpropagation code is correct, we usually disable gradient checking before using backpropagation to train the network.

Neural Networks: Learning

Quiz, 5 questions

3/5 points (60.00%)

Correct

Checking the gradient numerically is a debugging tool: it helps ensure a corre

ct implementation, but it is too slow to use as a method for actually computing gradients.

✖

0 / 1 point

5. Which of the following statements are true? Check all that apply.

- [ ] Suppose we have a correct implementation of backpropagation, and are training a neural network using gradient descent. Suppose we plot $J(\Theta)$ as a function of the number of iterations, and find that it is **increasing** rather than decreasing. One possible cause of this is that the learning rate α is too large.

This should be selected

- [ ] Suppose we are using gradient descent with learning rate α. For logistic regression and linear regression, $J(\theta)$

was a convex optimization problem and thus we did not want to choose a learning rate α that is too large. For a neural network however, $J(\Theta)$ may not be convex, and thus choosing a very large value of α can only speed up convergence.

Un-selected is correct

- [ ] If we are training a neural network using gradient descent, one reasonable "debugging" step to make sure it is working is to plot $J(\Theta)$ as a function of the number of iterations, and make sure it is decreasing (or at least non-increasing) after each iteration.

Correct
Since gradient descent uses the gradient to take a step toward parameters with lower cost (ie, lower $J(\Theta)$), the value of $J(\Theta)$ should be equal or less at each iteration if the gradient computation is correct and the learning rate is set properly.

- [ ] Suppose that the parameter $\Theta^{(1)}$ is a square matrix (meaning the number of rows equals the number of columns). If we replace $\Theta^{(1)}$ with its transpose $(\Theta^{(1)})^T$, then we have not changed the function that the network is computing.

Un-selected is correct