Machine Learning

Stanford University

About this Course

Machine learning is the science of getting computers to act without being explicitly programmed. In the past decade, machine learning has given us self-driving cars, practical speech recognition, effective web search, and a More vastly improved understanding of the human genome. Machine learning is so pervasive today that you probably use it dozens of times a day without knowing it. Many researchers also think it is the best way to make progress towns human-level Al. In this class, you will learn about the most effective machine learning techniques, and gain practice implementing them and getting them to work for yourself. More importantly banduland founding lead of theoretical underpinnings of learning, but also gain the practical know-how needed to quickly and powerfully apply these techniques to new problems. Finally, you'll learn about some of Silicon Valley's best practices in innovation as it pertains to machine learning and Al. English, Subtitles:

Chinese (Simplified), This course provides a broad introduction to machine Hebrew, Spanish, learning, datamining, and statistical half partern recognition. Topics include: (i) Supervised learning (parametric/non-parametric algorithms, support vectors and interest in machine learning). (ii) Unsupervised learning (plustering, dimensionality reduction, recommender systems, deep learning). (iii) Best practices in machine learning

(bias variance theory; innovation process in machine learning pass all graded (bias variance theory; innovation process in machine assignments to learning and Al). The course will also draw from complete the numerous case studies and applications, so that you'll also learn how to apply learning algorithms to building smart robots (perception, control), text understanding (web sear user Rajiags), computer is an analiging

informatics, audio, database mining அது முழ் என்ற

Syllabus

WEEK 1



Introduction

Welcome to Machine Learning! In this module, we introduce the core idea of teaching a computer to learn concepts using data—without being explicitly programmed. The Course Wiki is under construction. Please visit the resources tab for the most complete and up-to-date information.

- 5 videos, 9 readings expand
- 1. **Video:** Welcome to Machine Learning! **Graded:** Introduction
 - 2. **Reading:** Machine Learning Honor Code
- 3. **Video:** Welcome Linear Regression with One 4. **Video:** What is Machine Learning? Variable
- 5. **Reading:** What is Machine Learning? Linear regression predicts a real-valued output based **Grading** ut on the discussion application of linear regression to housing price prediction and introduce the gradient descent method reading. Supervised Learning
- 9. **Video:** Unsupervised Learning 7 videos, 8 readings expand
 - Reading: Unsupervised Learning
- 1. **Video:** Model Representation
 11**Gradiead ing** Whregres Mentwish One
 2.7. **Reading:** Model Representation
 - 2 Var**Reading:** Model Representation 12. **Reading:** Get to Know Your Classmates
 - 3. **Video:** Cost Function
 - 13. **Reading:** Frequently Asked Questions
- Lindea Reading: Cost Function Reading: Lecture Slides
- 5. **Video:** Cost Function Intuition I This optional module provides a refresher on

linear a**Reading**: @ Toption a linear algebra is necessary for the rest of the 7. **Video:** Cost Function - Intuition II course, especially as we begin to cover models with mare bingriables. The intuition II

- 96 vi**deo:** Gradient Descent E expand 109^{ui}**Reading:** Gradient Descent
 - 11. Video: Gradient Descent Intuition
 - 12. Reading: Gradient Answert Insuition
- 13: **Videe**: 原硬制硬件 和 实现有不 Linear **WEEK 2** Realtps idean
 - 14. **Reading:** And West And Secretary Linear Restantion
 Linear Regression with Multiple
 15. **Reading** at Rectain Multiple Variables
 - 6. **Reading:** Matrix Vector Multiplication What if your input has more than one value? In this modifie, what is Matrix Multiplication can be exterdeding. Matrix Multiplication features. We also discuss best practices for implementary Matrix Multiplication Properties
 - 10. **Reading:** Matrix Multiplication 8 yideos, 16 readings expand Properties
 - 11. **Reneding**: Versteings length of the Multiple 2Var Reneding: Inverse and Transpose
 - 2. **Reading:** Access MATLAB Online and 13. **Reading:** Lecture Slides Upload the Exercise Files

Otta Rearing Abit hip of all georg Windows

This.co. **Reading** destabling and airegons sign of the service and implendent the learning algorithms in practice.

To complete the programming assignments, 5. **Reading:** Installing Octave on Mac OS X you will need to use Octave or MATLAB. This (10.8 Mountain Lion and Earlier) module introduces Octave/Matlab and shows you how reading in the lines of the confidence of the confid

GNU/Linux

- 6 videos, 1 reading expand
 7. **Reading:** More Octave/MATLAB

- 9. Richerbingo Mipulitipher Ferra Durtes
- Video: Elloatding Dates cent for Multiple 10.

WEEK₅3 Variables Video: Control Statements: for, while, if

- 11. steading: Gradient Descent For Multiple
 - Variables **Video:** Vectorization

Logistic regression is a method for classifying character regression character regress might Eset logistic liegression to classify an email as spam or not spam. In this module, we 14. **Video:** Gradient Descent in Practice II introduce the notion of classification, the cost Learning Rate function for logistic regression, and the applicated in Bracilie at the same of the contraction of the contracti classificationing Rate

- 7 videos, 8 readings expand Regression
- 1. **Video:** Classification **Graded:** Logistic Regression 2. Reading Classification
- Video: Normalesquation Regularization 14: **Reading: NopotaleSquation**

Maghin vial pring models need to generalize lity well to new examples that the model has not seen in **Reading:** Nethern her being an Entereduce loninyertibilityhelps prevent models **Video:** Cost Function

- from overdee: Cost runction

 8. Reading: Cost Function
 Programming Assignments
 4 videos, 5 readings expand
 9. Video: Simplified Cost Function and
 23. Reading: Programming tips from
 - - 1. 例如如何中的Jem of Overfitting
- **Graded:** Regularizațion ading Eightigan Gest Loverium and Gradient Descent
 - Video: Cost Function
 - 11. Video: Advanced Optimization
 - **Reading:** Cost Function

WEEK24 **Reading:** Advanced Optimization

- **Video:** Regularized Linear Regression
- Video: Multiclass Classification: One-vs-
- 6. a **Reading:** Regularized Linear Regression

Nguraidectavonks: Representation

Nearal Readings Regularized hisposieticay how the

bration when your phone interprets and 9. Reading: Lecture Slides understand your voice commands, it is likely that a necogramming is helpting regressions tand your speech; when you cash a check, the machines that automatically read the digits also use neural networks.

- 7 videos, 6 readings expand
- - 2. **Video:** Neurons and the Brain
 - 3. Video: Model Representation I
- 4. **Reading:** Model Representation I **WEEK 5**
 - 5. Video: Model Representation II
 - 6. Reading: Model Representation II

Nouravide attendatives become interest

In this rReadings Examples and antuitions I backpronagation algorithm that is used to help learn parameters for a neural network. At the end of the discussifications in the end of the discussifications in the end of the discussification in the companies of the discussion in the companies of the companies of the discussion in the companies of the compan

- 28 vi**Readinge Mulitigs** aex palanssification
 - 13: **Reading** Lefture Slides
- Graded: Neural Networks: Learning 4: Reading: Classification and Neural Networks
 - 3. **Video:** Backpropagation Algorithm
- 4. **Reading:** Backpropagation Algorithm **WEEK 6**
 - 5. **Video:** Backpropagation Intuition
 - 6. **Reading:** Backpropagation Intuition

Advice ides: Applying Manbie enrolling Lear Pringeters

Applying fraction in practice is not always straightforward thris module, we share by the same term and discuss the best ways

to 1 evaluate a direct of the chief agned models.

- 17 villes, 7 Randoms Initialization
 - 12. **Reading:** Random Initialization1. **Video:** Deciding What to Try Next
- 3. Graded Apvison Fundation in Next 2. Learning: Evaluating a Hypothesis
 - 14. **Reading:** Putting It Together3. **Reading:** Evaluating a Hypothesis
 - 15. **Video:** Autonomous Driving

Machyideo: Model Selection and Design
16. Translateon restisses

To optimize a machine learning algorithm, you'll so received where the biggest read to first understand where the biggest frantive ideation/lest Sets improvements can be made. In this module, we discuss video Diageosing Diag vertoriance of a machine learning system with multiple parts, f. Reading: Diagnosing Bias vs. Variance and also how to deal with skewed data.

- 8. **Video:** Regularization and Bias/Variance
- 5 videos, 3 readings expand 9. **Reading:** Regularization and
- 1. Birited: in the Holding What to Work On Graded: Machine Learning System Design 12: Kiskning Priving What to Work On
 - 13. Reading: ropaknings Gurves
 - 14. **Kistening**e Eighne Wangs to Do Next

WEEK 7 Revisited

- 5. **Video:** Error Metrics for Skewed Classes
- 13. **Reading:** Deciding What to do Next 6. **Video:** Trading Off Precision and Recall

Supported to the Support Suppo

Suppor **Fregina** hang the stiffs Mest is negatione learning also sithmator slass with the We introduce the idea and intuitions behind SVMs and discuss how to use it in practice.

- 6 videos, 1 reading expand
- Video: Optimization Objective
 Graded: Support Vector Machines
 - 2. **Video:** Large Margin Intuition
 - 3. Video: Mathematics Behind Large

Margin Classification

WEEK₄8 Video: Kernels I

5. Video: Kernels II

- 6. Video: Using An SVM
- 7. **Reading:** Lecture Slides Unsupervised Learning
- 8. **Programming:** Support Vector We us Managerial Vised learning to build models that help us understand our data better. We discuss the k-Means algorithm for clustering that enable us to learn groupings of unlabeled data points.
- 5 videos, 1 reading expand
- 1. **Video:** Unsupervised Learning:
 - 2. Video: K-Means Algorithm

Dignerial Digner

In this nideo. Rendomoduitiel Zationpal
Components Analysis, and show how it can be 5. Video: Choosing the Number of Clusters used for data compression to speed up leaming and instructions of complex datasets.

- 7 videos, 1 reading expand
- - Video: Principal Component Analysis Problem Formulation

WEEK₄9 Video: Principal Component Analysis
Algorithm

5. **Video:** Reconstruction from Anomaly research Representation

Given a Video n Choosing that humber, whe may somethings want to move the average. For example, in manufacturing, we may want to detect defects or anor Reading a Gaussian distribution and how the model can be used for anomaly detection.

- 🗐 8 videos, 1 reading expand
- 1. **Video:** Problem Motivation **Graded:** Anomaly Detection

 2. **Video:** Gaussian Distribution
- 3. Video: AlgorithmRecommender Systems4. Video: Developing and Evaluating an
- When not beveraging and Evaluating an When not beverage when the products that automatically recommend other products that 5. Video: Anomaly Detection vs. you may like. Recommender systems look at Supervised Learning patterns of activities between different users

you may like. Recommender systems look at Supervised Learning patterns of activities between different users and difference productions. In this module, we 7. Video: Multivariate Gaussian

introduce recommender algorithms such as the Distribution collaborative filtering algorithm and low-rank materix Midenzationally Detection using the

Multivariate Gaussian Distribution

- 6 videos, 1 reading expand 9. **Reading:** Lecture Slides
- - Video: Content Based Recommendations
 - 3. Video: Collaborative Filtering

WEEK₄.10 Video: Collaborative Filtering Algorithm

5. **Video:** Vectorization: Low Rank Matrix Factorization

Large Scale Machine Learning
6. **Video:** Implementational Detail: Mean

Machine 使 种语模的 works best when there is an abundance of data to leverage for training. In this module, we discuss how to apply the machine regrammings: 他所以 中共安全的 and datasets.commender Systems

- 6 videos, 1 reading expand
- 1. Video: Learning With Large Datasets

 (2) Graded: Large Scale Machine Learning
 2. Video: Stochastic Gradient Descent
 - 3. Video: Mini-Batch Gradient Descent
- Video: Stochastic Gradient Descent
 WEEK 11_{Convergence}

5. Video: Online Learning

Application Example: Photo OCR Parallelism

Identifying and recognizing objects, words, and fights in an image is a challenging task. We discuss how a pipeline can be built to tackle this problem and how to analyze and improve the performance of such a system.

- 5 videos, 1 reading expand
- - 2. Video: Sliding Windows

View Less **Video:** Getting Lots of Data and Artificial

4. **Video:** Ceiling Analysis: What Part of the **How It Wooth S**o Work on Next

5. **Reading:** Lecture Slides

GENERAL Video: Summary and Thank You

How do I pass the course?

To earn your Course Certificate, you'll need to earn a passing grade on each of the required assignments—these can be quizzes, peer-graded assignments, or programming assignments. Videos, readings, and practice exercises are there to help you prepare for the graded assignments.

PROGRAMMING ASSIGNMENTS What do start dates and end dates mean?

Programming assignments require you if new rifes and run a computer program to solve a problem.

Once you enroll, you'll have access to all videos,

What are drivers, and programming assignments (if applicable). Peer-graded assignments can only be Proprietied in a reviewed to her best of the post of the least with the programment in the least reviewed to her best of the least reviewed to her best reviewed

assignments before the end of the session, you can

ethnell riestbuer næxt/sæsigionme/rotupraptogræssi millabeo isadisviblarad ovoiz/lobesatolesto Eairlik paprivibære i nglæleft off volum ginaskæstæsanobebægingsleted one at a time.

workafteravined attempted attempted by the property of the pro

witherpsecondankigheyouseschecktlerangdakterpyou coay seecortof wait piting ap inquizzes and programming assignified by the besubmitted late with our escive a gadelif you submit your peer-goaden assignified and programming assignification of the condition of th

Webs and the six fall reversing my the satisfall reversing my

Yes.

Eflya

grad

Structuring

MI Projects

you

Structuring
mprove your Arade, you can
yould neither two up to a graphe ent, we
jaie your mind your Provinces works a source or as

Focuses as many surrent ped of sold likely new foaye had a lasisan attestile mes ianvely to week for playing the loans as you programed to grasis lige for ment on quiz. We encourage

ie material during this delay. Convolutional Neural Networks

deeplearning.ai



Neural Networks and Deep Learning

deeplearning.ai



Sequence Models

deeplearning.ai